

MBrace: Cloud Computing with Monads

Jan Dzik Nick Palladinos Konstantinos Rontogiannis Eirik Tsarpalis Nikolaos Vathis

Nessos Information Technologies, SA
`fjdzik,npal,krontogiannis,eirik,nvathis@nessos.gr`

Abstract

pleted with examples and an informal description of semantics. We continue with distributed data (section 3), explaining the use and management of distributed data entities. In section 4 we present an overview of the MBrace runtime. Section 5 offers performance and scalability benchmarks on the MBrace framework, comparing results with Hadoop. In section 6, we give a brief overview of related work and finally, in section 7, discuss conclusions and future work.

2. Cloud Workflows

Cloud workflows form the essential pillar of MBrace; the program-

```
cloud {  
  for i in [ 1 .. 100 ] do
```

3. Distributed Data

Cloud workflows offer a programming model for distributed computation. But what happens when it comes to big data? While the distributable execution environments of MBrace do offer a limited form of data distribution, their scope is inherently local and almost certainly do not scale to the demands of modern big data applications. MBrace offers a plethora of mechanisms for managing data in a more global and massive scale. These provide an essential de-

Cloud Refs

The MBrace programming model offers access to persistable and distributed

```
let getRef () : Cloud<CloudRef<string []>> =  
  cloud {
```

```
    let! html = download "http://www.m-brace.org/0G0.40.40.4rg0.40.40.4RG[""]TJ0g0GJ/F317.9701Tf-20.922-8.967Td[(let!)]TJ/F2  
        rus, t(is)-774will,  
        in the theci tentin  
        computatio:d
```

MBrace comes out of the box with implementations for FileSystem, SQL and Azure storage providers, while providing pluggable,

References

- [1] Jeffrey Dean and Sanjay Ghemawat. MapReduce: simplified data processing on large clusters. In *Proceedings of the 6th conference on Symposium on Operating Systems Design & Implementation - Volume 6*